



Real-time Demo Workshop Rev. 3
(With additional exercises)


Table of Contents

Table of Contents	2
Workshop Setup Instructions	3
Creating a CrateDB Cluster.....	3
Creating AWS Workshop Pre-Requisites.....	4
Checking EC2 installation progress.....	10
View Grafana dashboards	11
Re-running the data producer.....	12
Change country for climate data	13
Additional Exercises	14
What is the time range?	14
Distribution of temperature	14
Average temperature per day	14
10 hottest points.....	15
Are we missing any data points?	16
Wind speed calculations	17
What are the top 10 windiest observations?	18
Wind direction calculated using trigonometry.....	18
Exercises - Geo.....	19
How to extract latitude & longitude from a GEO_POINT	19
Distance calculations	20
Exercises – Correlation	21
Is pressure and temperate related?.....	21
Is pressure and wind speed related?	21
Exercise – JSON object model	22

Workshop Setup Instructions

Creating a CrateDB Cluster

1. Go to <https://cratedb.com>
2. Click on the Start Free button
3. You have several options to create an account, use the one that is appropriate for you. You won't need to enter any payment details.



4. Once you have an account, log in and deploy a CRFREE cluster

Configure your new cluster

Shared
Single-node, suitable for non-production use cases with up to 8 shared vCPUs

Dedicated
Scalable up to 9 nodes, ideal for production workloads with up to 144 vCPUs

Custom
Customizable for large-scale needs, offers any cluster size and custom compute options

Shared tier clusters operate on a single node without replication, sharing vCPUs with other clusters. Performance may vary depending on the overall load, and usage is based on a fair-use principle. [Learn more about cluster types](#)

Cloud Provider

aws

A

G

Region

aws

AWS US-East (N.Virginia)
eks1.us-east-1.aws

[Request a new region](#)

Node compute size

CRFREE	Up to 2 vCPU	2 GiB RAM	FREE
S2	Up to 2 vCPU	2 GiB RAM	\$0.079 per hour
S4	Up to 3 vCPU	4 GiB RAM	\$0.158 per hour
S6	Up to 4 vCPU	6 GiB RAM	\$0.238 per hour
S12	Up to 8 vCPU	12 GiB RAM	\$0.475 per hour

Node storage size

8 GiB **FREE**

- Select a Shared cluster (this is the only way to get a free cluster)
- Select AWS as the cloud provider
- Select AWS US-East as the region (the same as the workshop runs within)

- Ensure CRFREE is selected as compute size
 - Click Deploy Cluster to create it, which will take a short amount of time
5. It's **very** important that you save the login credentials, so download or make a note now before moving on.
 6. We need to create the destination table for the climate data, run the following in the Console:

```
CREATE TABLE IF NOT EXISTS "demo"."climate_data" (  
  "timestamp" TIMESTAMP WITHOUT TIME ZONE,  
  "geo_location" GEO_POINT,  
  "data" OBJECT(DYNAMIC) AS (  
    "temperature" DOUBLE PRECISION,  
    "u10" DOUBLE PRECISION,  
    "v10" DOUBLE PRECISION,  
    "pressure" DOUBLE PRECISION,  
    "latitude" DOUBLE PRECISION,  
    "longitude" DOUBLE PRECISION  
  )  
);
```


7. The host information you will need later, so make a note or leave the webpage open for now.

Creating AWS Workshop Pre-Requisites

Note: Please use the exact names provided in this guide for naming instances otherwise build steps might fail. Use copy paste to avoid errors.


Note: If you already have AWS console it's best to use another browser

1. Click on the provided link and add the email you registered for the event to get an OTP
2. After entering the OTP, you received go to the Amazon Console




The screenshot shows the AWS Workshop Studio interface. On the left, there's a sidebar with navigation links like 'Introduction', 'Lab Setup', 'Source Connectors', 'Sink Connectors', 'Clean Up', 'Known Issues', and 'Authors'. Below this, there's a section for 'AWS account access' with a red box highlighting the 'Open AWS console (us-east-1)' link. The main content area displays event information for 'CrateDB in Munich 27/11 test event', including start time, duration, accessible regions, and a description. It also shows workshop details like title, complexity level, topics, and AWS services. At the bottom, there's a section for 'Event outputs' which is currently empty.

3. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>
Choose region us-east-1 this will be the region that you are using for the workshop.




The screenshot shows the AWS Management Console. The top navigation bar indicates the region is 'us-east-1'. The left sidebar shows the 'EC2' service selected. The main content area displays the 'Resources' section for EC2, showing a table of instances, key pairs, and other resources. The 'Launch instance' button is highlighted. The right sidebar shows the 'Account attributes' section, including the account name and various settings.

2. Choose Key Pairs in the navigation bar on the left and click on Create key pair.



The screenshot shows the AWS Management Console with the 'Key Pairs' page selected in the left sidebar. The main content area displays a table of key pairs. The 'Create key pair' button is highlighted in the top right corner of the page.

- Provide a name for the key pair (e.g. mskconnect-workshop-keypair) and click on Create key pair.



Create key pair [Info](#)

Key pair
A key pair, consisting of a private key and a public key, is a set of security credentials that you use to prove your identity when connecting to an instance.

Name

The name can include up to 255 ASCII characters. It can't include leading or trailing spaces.


Key pair type [Info](#)
☒ RSA
☐ ED25519

Private key file format
☒ .pem
 For use with OpenSSH
☐ .ppk
 For use with PuTTY

Tags (Optional)
No tags associated with the resource.
[Add tag](#)
 You can add 50 more tags.


[Cancel](#) [Create key pair](#)

- You should see the message Successfully created key pair, and confirm the download of the generated .pem file to your local machine.




- Download the CloudFormation template here:
<https://github.com/crate/realtime-demo/releases/download/CloudFormation/CloudFormationTemplateWithLambdaV2.yaml>
- Open the AWS CloudFormation console at
<https://console.aws.amazon.com/cloudformation/>

7. Click on Create stack and select With new resources (standard).




8. Choose Upload a template file option and upload the CloudFormation template from your local machine, then click Next.



9. Enter mskconnect-workshop as the stack name. There are parameters on the next screen that need to be entered correctly, although some have sensible default values (and likely won't need to be changed). Each parameter has a description to help identify what it needs to be configured as but ask for help if unsure.

Choose mskconnect-workshop-keypair under KeyName.

Then click Next.



- For the step Configure stack options, ensure Delete newly created resources during a rollback is set to Delete all newly created resources.

Stack failure options

Behaviour on provisioning failure

Specify the roll-back behaviour for a stack failure. [Learn more](#)


- ☒ **Roll back all stack resources**
Roll back the stack to the last known stable state.
- ☐ **Preserve successfully provisioned resources**
Preserves the state of successfully provisioned resources, while rolling back failed resources to the last known stable state. Resources without a last known stable state will be deleted upon the next stack operation.

Delete newly created resources during a rollback

Specify whether resources that were created during a failed operation should be deleted regardless of their deletion policy. [Learn more](#)

- ☐ **Use deletion policy**
Retains or deletes created resources according to their attached deletion policy.
- ☒ **Delete all newly created resources**
Deletes created resources during a rollback regardless of their attached deletion policy.


Scroll down to the bottom and check the box to acknowledge the creation of IAM resources.



The screenshot shows the AWS CloudFormation console interface. The 'Configure stack options' page is displayed, showing various configuration sections. The 'Capabilities' section at the bottom contains a checkbox labeled 'I acknowledge that AWS CloudFormation might create IAM resources.' which is checked and highlighted with a red box. The 'Create stack' button at the bottom right is also highlighted with a red box.

Click Next and then Submit to trigger the creation of the CloudFormation stack.

- Wait until the stack creation is completed before proceeding to the next steps. It may take a while for the MSK cluster creation (~30 to 45 mins).



What does the above CloudFormation template do?

The CloudFormation stack creates the following resources:


- 2 CloudWatch Log Groups for MSK Cluster and MSK Connect connectors respectively, each with 7 days log retention period
- 1 S3 bucket
- 1 VPC with 1 Public subnet and 3 Private subnets, and the required networking components such as Route Tables, Internet Gateway, NAT Gateway & Elastic IP, VPC Gateway Endpoint for S3
- 1 m5.large EC2 instance with the required Security Group, IAM Role, EC2 Instance Profile. On this instance various things will be installed including Python 3.13, netCDF4 and other libraries needed to execute the data producer.
- 1 MSK Cluster with 3 kafka.m5.large broker nodes configured with Apache Kafka version 3.6.0, and the required Security Group

Note that the encryption options are disabled to reduce the chance of errors in the workshop, this is not a best practice, and you should not use the provided CloudFormation template in any production environments.

Checking EC2 installation progress


Open the EC2 dashboard <https://console.aws.amazon.com/ec2/> and go to Instances

Wait for the instance (named CrateDBInstance) to be fully ready before attempting to connect (it will show state as running).



The screenshot shows the AWS Management Console 'Instances' page. The instance 'CrateDBInstance' is listed with ID 'i-0c283a8cdf682c9bc' and is in the 'Running' state. The instance type is 'm5.xlarge' and the status check shows '3/3 checks passed'.

Select it and click on the Connect button.



The screenshot shows the 'Connect to instance' page in the AWS Management Console. The 'EC2 Instance Connect' tab is selected. The instance ID 'i-0c283a8cdf682c9bc' is shown. The 'Connection type' is set to 'Connect using a Public IP'. The 'Public IPv4 address' is '44.210.127.196'. The 'Username' is 'ec2-user'. A note states: 'In most cases, the default username, ec2-user, is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI username.'

In the shell use the following command to follow progress:

```
tail -f /var/log/user-data.log
```


Wait until the following is displayed (note this may take some time):

Everything is installed!

View Grafana dashboards

Your dashboard should be accessible at <http://ip:3000> (where IP is the assigned external IP address of the EC2 instance).

To login, the default Grafana username and passwords are admin/admin, then either skip or set a new password.



All complete!

Re-running the data producer

The data producer will run during the installation, but if you want to run the demo again and want to see the visualization in real time building up, simply go back to your query console and truncate the table:

```
delete from demo.climate_data;
```

Then run the following from the EC2 instance:

```
cd ~/realtime-demo/data  
source .venv/bin/activate  
python3 producer.py
```

Change country for climate data

If you like, you can modify the source data to be local to the Netherlands. It will require to get an account for the weather data and the data to be downloaded by the producer code which can take some time. Here are the steps of how to do it.

Change Country to Netherlands (NLD)

In your parser code initialization, change:

```
parser = ClimateParser("DEU")
```

to:

```
parser = ClimateParser("NLD")
```

NOTE: Clear Old Data (delete table) BEFORE starting the producer

Delete the existing data so the dashboard shows only fresh Netherlands data:

```
delete from demo.climate_data;
```

Start the Producer Again

Now run the producer to download and ingest new data:

```
python3 producer.py
```

The pipeline retrieves ERA5-Land data for the Netherlands, parses it, and produces JSON records.

View Updated Grafana Dashboard

Open Grafana and refresh the dashboard.

It will now display only the newly ingested Netherlands data.

Additional Exercises

Try some queries to see the flexibility in the data model

What is the time range?

```
SELECT
  min("timestamp") AS start_ts,
  max("timestamp") AS end_ts
FROM
  demo.climate_data;
```

Timestamps in CrateDB are always stored without a timezone. If the field is defined WITH a timezone, then on ingest the timestamp is adjusted accordingly to UTC.

In this data, it's defined without a timezone. It's defined as the number of milliseconds since the Unix epoch (`1970-01-01T00:00:00Z`). For a detailed guide to all the available datatypes see:

<https://cratedb.com/docs/crate/reference/en/latest/general/ddl/data-types.html#timestamp>

Note: The CrateDB console will display both the number of milliseconds and the human readable date and time.

Distribution of temperature

```
SELECT
  min(data['temperature']) AS min_t,
  avg(data['temperature']) AS avg_t,
  max(data['temperature']) AS max_t
FROM
  demo.climate_data;
```

This is a straightforward aggregation query. It's running a query across all data in the table and returning the minimum, maximum and average temperatures.

Notice the **AS**, which is an alias. Normally fields in an aggregate will be unnamed expressions, but with recent versions of CrateDB it will generate an alias automatically. However, it's still recommended to give a name.

Average temperature per day

```
SELECT
  date_trunc('day', "timestamp") AS day,
  min(data['temperature']) AS min_temp,
  max(data['temperature']) AS max_temp
FROM
  demo.climate_data
GROUP BY
  1
ORDER BY
  1;
```

This query introduces a new function - `date_trunc`. This returns a timestamp value *truncated* to a specific precision, in this case to precision of a day and effectively ignores the HH:MM:SS of any timestamp. This gives us a per-day aggregate figure. To get information on more functions like this, see:

<https://cratedb.com/docs/crate/reference/en/latest/general/builtins/scalar-functions.html>

10 hottest points

```
SELECT
  "timestamp",
  geo_location,
  data['temperature'] AS temp
FROM
  demo.climate_data
ORDER BY
  temp DESC
LIMIT
  10;
```

To give us the top 10 hottest points, we use `ORDER BY` (to order by temperature in descending order) and `LIMIT` to only return 10 rows.

The `SELECT` part of the query returns us the timestamp (i.e. when), the location (as a `GEO_POINT` co-ordinate, which is explained later) and the temperature from the `data` JSON object.

Are we missing any data points?

```
SELECT
  count(*) FILTER (WHERE data['temperature'] IS NULL) AS missing_temp,
  count(*) FILTER (WHERE data['pressure'] IS NULL) AS missing_pressure,
  count(*) FILTER (WHERE geo_location IS NULL) AS missing_geo
FROM
  demo.climate_data;
```

As you can see, nothing is missing in the included data. How does this query work?

FILTER is a window function. It performs a calculation across a set of related rows defined by an expression (or multiple expressions). In this case we are **COUNT**ing the number of rows where one of `data['temperature']`/`data['pressure']` or `geo_location` is NULL (i.e. undefined).

Window functions are very powerful, and it's very much recommended to read the following webpage to learn more about their capabilities.

<https://cratedb.com/docs/crate/reference/en/latest/general/builtins/window-functions.html>

Let's add a point with a missing temperature and re-run:

```
INSERT INTO demo.climate_data (timestamp, geo_location, data) VALUES
(1754956800000, [8.788311111111111, 54.903], {"longitude" = 8.788311111111111,
"latitude" = 54.903, "u10" = 4.472952365875244, "v10" = -1.3958832025527954,
"pressure" = 102426.1015625});
```

If you now re-run the missing data point query, it will show:

```
missing_temp missing_pressure missing_geo
1           0                0
```

Now let's update the record with a valid temperature:

```
UPDATE demo.climate_data
SET
  data['temperature'] = 16.8
WHERE
  timestamp = 1754956800000
  AND geo_location = [8.788311111111111, 54.903];
```

Re-run the missing data point query, and you'll see it now shows zero again.

Wind speed calculations

This query will calculate min/max and avg wind speeds (in metres per second) using Pythagoras' theorem to calculate the diagonal (hypotenuse) value:

$$\sqrt{(x^2 + y^2)}$$

```
SELECT
  min(sqrt(power(data['u10'], 2) + power(data['v10'], 2))) AS min_ws,
  avg(sqrt(power(data['u10'], 2) + power(data['v10'], 2))) AS avg_ws,
  max(sqrt(power(data['u10'], 2) + power(data['v10'], 2))) AS max_ws
FROM demo.climate_data;
```

The **power** function returns the field to the specified power, 2 in this case, i.e. it squares the field value. The **sqrt** functions returns the square root. These are part of the CrateDB scalar functions, and more information is available here:

<https://cratedb.com/docs/crate/reference/en/latest/general/builtins/scalar-functions.html>

Can we optimise this query? The mathematical expression is being calculated 3 times and this is inefficient. We can use a CTE (common table expression) to optimise this:

```
WITH wind_speed AS (
  SELECT
    sqrt(power(data['u10'], 2) + power(data['v10'], 2)) AS ws
  FROM demo.climate_data
)
SELECT
  MIN(ws) AS min_ws,
  AVG(ws) AS avg_ws,
  MAX(ws) AS max_ws
FROM wind_speed;
```

Now we are only calculating once. For more information on common table expressions, see:

<https://cratedb.com/querying/common-table-expressions-ctes>

What are the top 10 windiest observations?

```
SELECT
  "timestamp",
  geo_location,
  sqrt(power(data['u10'], 2) + power(data['v10'], 2)) AS wind_speed
FROM
  demo.climate_data
ORDER BY
  wind_speed DESC
LIMIT 10;
```

This uses the wind speed calculation as detailed in the previous query.

Wind direction calculated using trigonometry

This query will show the wind direction in degrees.

```
SELECT
  date_trunc('hour', "timestamp") AS hour,
  avg(sqrt(power(data['u10'], 2) + power(data['v10'], 2))) AS avg_speed,
  avg(
    mod(degrees(atan2(data['u10'], data['v10']))) + 360, 360)
  ) AS avg_dir_deg
FROM demo.climate_data
GROUP BY 1
ORDER BY 1;
```

The function `mod` has been introduced here, this refers to the **modulus** operation (`mod` is a shortened alias to **modulus**). This function returns the remainder of a divide operation; in this case we use it so that the angle in degrees is always between 0 and 360, i.e. an angle of 370 would return 10.

The other new function is `atan2`, which is a trigonometry function. All these are covered in the scalar functions guide:

<https://cratedb.com/docs/crate/reference/en/latest/general/builtins/scalar-functions.html#scalar-modulus>

Exercise: try returning the result in radians (ensure you update the modulo operation, the `pi` function will be useful here!)

Exercises - Geo

How to extract latitude & longitude from a GEO_POINT

You can use the `latitude` and `longitude` functions to return the co-ordinates in a `GEO_POINT` datatype. This example will return the bounding box for all loaded data:

```
SELECT
  min(longitude(geo_location)) AS min_longitude,
  max(longitude(geo_location)) AS max_longitude,
  min(latitude(geo_location)) AS min_latitude,
  max(latitude(geo_location)) AS max_latitude
FROM
  demo.climate_data;
```

Let's use the result of this to work out how many data points there are in a subset of the data, take this query and fill in some correct values based on the result of the above query.

```
SELECT count(*) AS points
FROM demo.climate_data
WHERE latitude(geo_location) BETWEEN 0.0 AND 0.0
      AND longitude(geo_location) BETWEEN 0.0 AND 0.0;
```

If you get it right, this should return a non-zero count. If it's zero, check your query and try again.

More information on CrateDB's geospatial capabilities is detailed here:

<https://cratedb.com/docs/guide/start/modelling/geospatial.html>

A more powerful way of doing this is to use the `within` function. This allows us to find points that lie within a polygon, which gives us the flexibility to use shapes other than a rectangle:

```
SELECT count(*) AS points
FROM demo.climate_data
WHERE within(
  geo_location,
  {
    "type": "Polygon",
    "coordinates": [[
      [0.0, 0.0],
      [0.0, 1.0],
      [1.0, 1.0],
      [1.0, 0.0],
      [0.0, 0.0]
    ]]
  }
);
```

The list of points is an array, in the form:

```
[longitude1, latitude1],  
[longitude2, latitude2] ,  
[longitude3, latitude3] ,  
[longitude4, latitude4] ,  
[longitude1, latitude1]
```

Which effectively *defines* a rectangular bounding box.

Note: Don't be afraid to ask for some guidance from one of the CrateDB team.

Distance calculations

CrateDB also supports distance calculations, the following query gives the nearest points to Berlin, Germany (approximate location is 52.5N, 13.4E)

```
SELECT  
  geo_location,  
  distance(geo_location, 'POINT(13.4 52.5)') AS metres,  
  data['temperature'] AS temp  
FROM demo.climate_data  
GROUP BY geo_location, metres, temp  
ORDER BY metres ASC  
LIMIT 50;
```

Exercise: Try changing the location from Berlin to some other location.

Exercises – Correlation

Is pressure and temperate related?

Is there a relationship between temperature and atmospheric pressure (at ground level)?

Run this query and decide for yourself:

```
SELECT
  geo_location,
  avg(data['temperature']) AS avg_temp,
  avg(data['pressure']) AS avg_pressure
FROM demo.climate_data
GROUP BY 1
ORDER BY 2 DESC;
```

Exercise: Is there a way to improve this query? Maybe a combination of temperature variance from average and air pressure from 100,000 (1 atmosphere)?

Is pressure and wind speed related?

Is there a relationship between temperature and wind speed?

```
SELECT
  geo_location,
  avg(data['temperature']) AS avg_temp,
  avg(data['pressure']) AS avg_pressure,
  avg(sqrt(power(data['u10'], 2) + power(data['v10'], 2))) AS wind_speed
FROM demo.climate_data
GROUP BY 1
ORDER BY 3 DESC;
```

Exercise: Try altering the ordering, either **ASC**ending or **DESC**ending, or using a different column (or combination).

Exercise – JSON object model

Insert a new value into the dataset.

```
INSERT INTO demo.climate_data (timestamp, geo_location, data) VALUES
(123, [8.788311111111111, 54.903], {"longitude" = 8.788311111111111,
"latitude" = 54.903, "temperature" = 16.868310546875023, "u10" =
4.472952365875244, "v10" = -1.3958832025527954, "pressure" =
102426.1015625});
```

This is using the existing schema. What if we want to add a new humidity value to our dataset, do we need to recreate the table?

In CrateDB we can define object columns (i.e. JSON) as one of three different types:

OBJECT(STRICT) – The schema for the object is fixed

OBJECT(DYNAMIC) – The schema is dynamic, and new attributes can be added

OBJECT(IGNORED) – The JSON is stored ignoring the datatypes of the attributes and not indexing the field

<https://cratedb.com/blog/handling-dynamic-objects-in-cratedb>

When you created this table you set the **data** column to be DYNAMIC, so we can add additional attributes easily, and these attributes will be indexed.

Add a new key in the JSON object without the need to modify the schema

```
INSERT INTO demo.climate_data (timestamp, geo_location, data) VALUES
(123, [8.788311111111111, 54.903], {"humidity" = 43.4, "longitude" =
8.788311111111111, "latitude" = 54.903, "temperature" =
16.868310546875023, "u10" = 4.472952365875244, "v10" = -
1.3958832025527954, "pressure" = 102426.1015625});
```

Now try to query the new key:

```
SELECT
  data['humidity']
FROM
  demo.climate_data
WHERE
  data['humidity'] IS NOT NULL;
```

This makes it very easy to extend the data model, ideal for IoT applications where additional sensor information may become available.

And of course, with CrateDB all of this additional data will be indexed for great query performance.

Exercise: Try adding additional columns of data or more data and add this view to the Grafana dashboards.